

# Чифир (Chifir) Engine

Copyright 2025 Randomcode Developers

## Contents

1. Introduction .....	2
1.1. Language .....	2
1.2. Target platforms .....	2
1.3. Benefits for developers .....	2
1.4. Design .....	2
1.5. Systems .....	2
2. Engine components .....	3
3. Platforms .....	5
3.1. x86 .....	5
3.2. Windows .....	5
3.3. Linux .....	6
4. Libraries .....	7
5. Engine coding style .....	8
5.1. Colourful language .....	8
5.2. Documentation .....	8
5.3. External code .....	8
5.4. File header .....	8
5.5. Formatting .....	8
5.6. #includes .....	8
5.7. Types .....	9
5.8. Comments .....	9
5.9. Naming .....	9
5.10. Classes vs structs .....	9
5.11. Headers .....	9
5.12. Standard library replacement .....	9
5.13. Assertions and error handling .....	10
6. Tools .....	11
7. Scene system .....	12
8. Renderer architecture .....	13
8.1. Hardware interface .....	13
8.2. Rendering pipeline .....	13
8.3. Render system .....	13

# **1. Introduction**

This document outlines the design of the custom-made 3D game engine for False King and subsequent Randomcode Developers games.

## **1.1. Language**

The engine is being written in C++14, with a custom runtime/container library and almost no libraries, maximizing portability.

## **1.2. Target platforms**

The engine is designed to be extremely portable. It currently supports Windows and Linux, and will target Xbox Series X|S, PlayStation 5, and Nintendo Switch/Switch 2.

## **1.3. Benefits for developers**

The engine will be as open source as possible (essentially, anything that I'm legally allowed to open source). It will give me and any interested developers a solid foundation to build advanced games on. It also benefits me as a huge learning experience.

## **1.4. Design**

The engine will be based on an entity component system, fairly clean separation between independant components with certain common ones, like platform abstraction, and simple data formats.

## **1.5. Systems**

Each system will be a static or dynamically loaded shared library, and expose a general interface in addition to ECS systems. This will make it easy to add and integrate new features.

## 2. Engine components

The engine will be made of these pieces:

Component	Components needed	Functionality	Available in tools builds
Base	system calls, libc on non-Windows	containers, basic algorithms, strings, data manipulation and serialization, Unicode handling, startup, shutdown, threading, synchronization, screen output, system information, basic file system functions, input, debugging features, logging	yes
VideoSystem	Base	abstracts a window or game console screen	yes
Rhi*	Base, VideoSystem	abstracts Vulkan, DirectX 12, GNM, OpenGL, DirectX 8/9, whatever else	maybe?
Rhi	Base, Rhi*, VideoSystem	exposes RHI backends	maybe?
Math	none	implements linear algebra and other math	yes
Texture	Base	texture format	yes
Mesh	Base	mesh format	yes
Pack	Base	package file format	yes
LauncherMain	none	loads Launcher	yes
Launcher	Base	loading an application DLL and the components it needs	yes
Engine	Base	cameras, scene management, entity component system, commonly used components (for entities), system management	no
RenderSystem	Base, Math, Rhi, VideoSystem	rendering scenes, UIs, anything else	no
InputSystem	Base	user input	no
UiSystem	Base, InputSystem, Math, RenderSystem	user interfaces	no

PhysicsSystem	Base, Math	simulates mechanical physics	no
AnimationSystem	Base, Math	controls skeletal animation	no
AudioSystem	Base, Math	handles audio	no
Game	AnimationSystem, Base, Engine	game functionality common between client and server, such as prediction and data parsing	no
GameServer	Base, Engine, PhysicsSystem	game functionality that happens on the server, such as simulation, player management, etc	no
GameClient	Base, Engine, InputSystem, RenderSystem, UiSystem	game functionality that happens on the client, such as rendering, player input, and possibly prediction	no

### 3. Platforms

The engine will support at least Windows and Linux. All desktop platforms will use Steam, all others will use the platform's official store.

Platform	Toolchain	Graphics API(s)
<b>Current progress</b>	Windows	MSVC, GDK
DirectX 12, Vulkan, OpenGL	Most complete port	Linux
LLVM	Vulkan, OpenGL	Largely complete
Xbox Series X S	MSVC, GDKX	DirectX 12
Minimal effort with existing Windows support	PlayStation 4	LLVM, PS4 SDK
GNM	Not too hard, other than GNM	PlayStation 5
LLVM, PS5 SDK	GNM	See PS4 note
Nintendo Switch/Switch 2	LLVM, Switch SDK	Vulkan
Not supported yet		

These platforms may be supported purely out of personal interest:

Platform	Toolchain	Graphics API(s)	Notes
Xbox 360	Ancient MSVC	DirectX 9	The engine builds, but somehow the XEXs don't have export tables, meaning it doesn't run
PlayStation 3	Ancient GCC, possibly modern LLVM	GCM, OpenGL	Haven't tried this very hard yet, it's probably possible
PlayStation Portable	GCC	OpenGL	Crashes in homebrew startup code, seems like current firmware doesn't support how syscalls are used in it
Bare metal x86	LLVM	Software renderer	This would take a lot of engineering and probably not be worth it

#### 3.1. x86

On x86-based platforms, SIMD is dynamically detected at startup using `cputid`, so old CPUs still work but modern ones can be used fully. Unfortunately, this does reduce the opportunities for optimization on 32-bit, but that's fine because that isn't going to be the main build that people get. On consoles, the exact `CPUID` results could be hardcoded.

#### 3.2. Windows

On Windows, the engine does some pretty crazy things. For one, it directly uses `ntdll.dll` instead of `kernel32.dll`, which is mostly a matter of preference. Another thing is that `Base.dll` manually imports system functions from `ntdll.dll` and `user32.dll` and exports them for other modules, in addition to exporting `*_Available` functions. This all lets it run on ancient versions, but then dynamically importing useful functions from newer versions when they're available. The manual

importing works by having function pointers and exporting forwarder functions that call them as the real names that `ntdll.dll` or whatever else exports, and then also having functions that check whether the function pointers are null or not. It even avoids having an import table at all by using the PEB to get `ntdll.dll`'s base address, parsing it, and finding `LdrGetProcedureAddress`, and then using that to get other functions normally.

Additionally, UWP is supported dynamically as well. If the engine detects that it's running in a packaged context, then it uses WinRT through COM interfaces instead of Win32 for windowing. It still mostly uses functions from `ntdll.dll` otherwise, though.

### **3.3. Linux**

Linux is supported, but relies on `libc` for timezones, startup, and library loading.

## 4. Libraries

Library	Use
<u>phnt</u>	Exposes internal Windows APIs
<u>stb</u>	Custom <code>snprintf</code> and other functionality
<u>volk</u>	Vulkan loader
<u>Vulkan-Headers</u>	Used for accessing Vulkan APIs
<u>musl</u>	Some code borrowed for handling time on Linux

## 5. Engine coding style

The engine is written in C++17 with no STL and (almost) no C runtime. This comes with some advantages, but plenty of disadvantages as well.

Pros	Cons
Full control over nearly everything	Have to implement everything
Very portable, almost everything that can possibly run the engine has a C++14 compiler	The Xbox 360 is a platform I'd really like to support, but it only has a C++03 TR1/C89 compiler
Only the exact functionality needed is implemented	Features not built on existing support code take longer

Just try to copy the existing style as much as possible and you'll be fine. Don't reformat external code.

### 5.1. Colourful language

You shouldn't be mean to anyone, but you can swear a little and be snarky. You can complain about something that gave you a headache within reason.

### 5.2. Documentation

You should try to document everything important. Any public declarations, macros, constexpr functions, any weird decisions, folders with weird stuff in them, anything where someone might wonder what it is and how it works/why it's there. For implementations of interfaces/abstract classes, you don't have to document inherited functions. TLDR, just document stuff in a header file and anything complicated/weird.

### 5.3. External code

Always document where external code is from, and avoid introducing it. Put it in external when possible, preferably as a Git submodule unless it's only a few files. Always include licenses, and add information to `scripts/licenses.toml`.

### 5.4. File header

Put a very short comment followed by a copyright line (usually for Randomcode Developers). Keep track of where files outside external come from, and add a license in `external` and `scripts/licenses.toml` for loose code (i.e. SDL/musl).

Example:

```
++  
/// \file <file description>  
/// \author Randomcode Developers
```

### 5.5. Formatting

Just use `clang-format` aggressively. The only thing that I'm pretty sure it doesn't do is adding curly brackets to one-line if statements and loops, which is part of the style. Also, be careful about putting a blank line between headers that shouldn't be sorted alphabetically.

### 5.6. #includes

Separate by folder, sort alphabetically when possible



## 5.7. Types

`public/base/types.h` defines short type names largely based on Rust's type names. Sizes should use the signed `ssize` to make calculation errors easier to see, and the fact that it's a size already gives the indication it can't be negative. Any other type can be unsigned, and for certain things like `operator new()` where using `usize` is required, that's fine too.

## 5.8. Comments

Comments should explain what code does. At the top of a file that implements something complex, or the main header for a whole component, explain the overall design of the component, any important choices and the reasoning, and whatever limitations exist. Additionally, when functions are complex, add comments explaining what's happening/why it's happening. The memory manager in `base/memory.cpp` and the Windows loader in `base/loader_win32.cpp` are the best example of commenting things so far.

## 5.9. Naming

Variables are camel case, prefixed with `m_` for private/protected members, `g_` for globals, and `s_` for static globals, and `f_` for (some) function pointers. Types are Pascal case, prefixed with `C` for classes, `CBase` for abstract classes, and `I` for interfaces, and typedefs end with `_t`. Functions are Pascal case, with the name of their component and an underscore as a prefix, like `Base_`. Common abbreviations (like `str`, `len`, `max`, `min`, `alloc`, `buf`, `src`, `dest`, common acronyms) are allowed, but obscure ones should be avoided. Try to balance clearness with succinctness when naming variables, so they're easier to type but you can still easily recognize them.

## 5.10. Classes vs structs

Classes do things, structs store data (you can have a destructor in them though). That's the distinction so far.

## 5.11. Headers

Public headers (ones visible to any component) should include as few headers as possible, and forward declare types where needed. In `.cpp` files, all headers for the things used should be included, not just ones that happen to include the right things. Private headers are more free to include things, and have references to globals inside components, like `base/base.h`.

## 5.12. Standard library replacement

Because the C runtime and STL aren't used, there are some replacements for the commonly used stuff, and there are also utility functions commonly implemented on top of these, like automatically allocating a buffer for `snprintf`.

In terms of replacements for the CRT, `public/base/base.h` has `Base_Alloc`, `Base_MemSet`, `Base_MemCopy`, and `Base_MemCompare`, and `public/base/basicstr.h` has `Base_StrFormat`, `Base_StrCopy`, `Base_StrClone`, and `Base_StrCompare`. They work basically just like `malloc`, `memset`, `memcpy/memmove`, `memcmp`, `snprintf`, `strcpy`, `strdup`, and `strcmp`, but because this is still C++, they're overloaded and have behaviour controlled by parameters, which makes them more convenient to use. `Base_MemSet`, `Base_MemCopy`, and `Base_MemCompare` (and the string functions implemented on top of them) also make use of SIMD where possible.

There's not many fancy containers yet, but `CVector<T>` defined in `public/base/vector.h` is a working implementation of a dynamic array. `public/base/string.h` defines `CString`, and it implements advanced features like splitting and multiplication. Additionally, there's `CIntrusiveLinkedList<T>`, which is used for the free list in the memory allocator, and offers significant user control over the nodes for exactly that reason.

### 5.13. Assertions and error handling

Assertions are mainly for scenarios that shouldn't happen, and are disabled in retail builds because anything triggering them should be caught in debug/release builds; don't use them for general error handling. For example, if a piece of memory *must* be allocated successfully, like in `operator new()` where the standard technically requires that it not return `nullptr` (even though the standard isn't as relevant for the engine), or an index is outside the valid range, or a parameter is wrong in a way it shouldn't be, then you can use an assert. Normally, you can use the `ASSERT` macro. If a condition isn't the most indicative of why something is wrong, `ASSERT_MSG` lets you add a message. For functions which just succeed or fail, return `false`, `nullptr`, or some other reasonable/documented value when an error happens. When an unrecoverable error happens, use `Base_Quit` (or `Base_Abort`/`Base_AbortSafe` in functions where logging/allocation isn't available) to kill the engine and show the user an error message.

## 6. Tools

Tool	Use	Custom made?
<u>xmake</u>	Build system	no
<u>Visual Studio 2022</u>	Code editing, debugging (Windows, consoles)	no
<u>Visual Studio Code/Neovim</u>	Code editing (non-Windows)	no
<u>GDB/LLDB</u>	Debugging (non-Windows)	no
<u>DXC</u>	Shader compiler	no
<u>spirv-cross</u>	Shader converter	no

## 7. Scene system

Scenes contain entities. Visible entities have components like a mesh or a camera, and transform information. These are some categories of entities:

- Sky/sun/moon entities
- Details like grass, leaves, etc
- Terrain
- Objects like furniture, items, etc
- Players
- NPCs
- Buildings, doors, etc

## 8. Renderer architecture

The renderer will be implemented in multiple layers, flexible enough to support drawing and post-processing fairly complex scenes, extensible with more techniques and passes, and simple to use for the rest of the engine.

### 8.1. Hardware interface

The hardware interface is an abstraction of Vulkan/Direct3D/GNM/whatever other ungodly API I have to deal with. It's low level, and implements render targets, materials, and geometry primitives, as well as special render targets just for going to the screen (they wrap the swap chain images).

- Handles `VkInstance/IDXGIFactory`, `VkDevice/ID3D12Device`, `VkCommandBuffer/ID3D12GraphicsCommandList`, `VkSwapChainKHR/IDXGISwapChain`
- Creates and manages geometry (VB+IB), textures, render targets, shaders, materials (texture + shader)
- Handles drawing given geometry + material

### 8.2. Rendering pipeline

Handles the process of taking data (model, position, etc of objects in scene, and general properties of the world) and using the hardware interface to render and post-process all of it.

- Calls for drawing objects and adding lights
- Uses multiple render passes to light and post-process the scene
- Rasterization-based lighting passes
- Ray-tracing-based passes
- Common post-processing passes

### 8.3. Render system

Calls into the rendering pipeline to draw scenes from different cameras, such as the player's eyes/over the shoulder, cinematic cameras, mirrors and other reflective surfaces, and literal cameras.

- ECS system that iterates over objects in the scene
- Sets parameters based on scene, such as sky details (even that could be an entity)

